


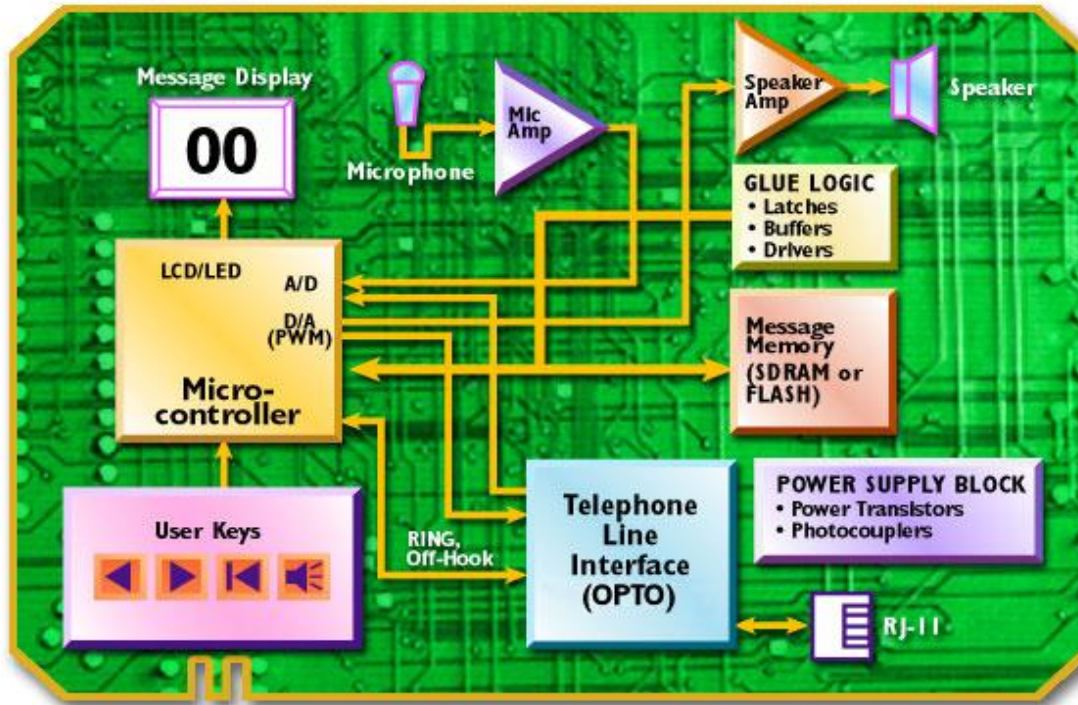
Preemptive Scheduling

These lecture notes created by Mr. James B. (Jim) Carlson, NCSU

Big Picture

- Methods learned so far
 - We've been using a *foreground/background* system
 - Interrupt service routines run in foreground
 - Task code runs in background
 - Limitations
 - Must structure task functions to run to completion, regardless of “natural program structure” – can only yield processor at end of task
 - Response time of task code is not easily controlled, in worst case depends on how long each other task takes to run
- What we will learn next
 - How to share processor flexibly among multiple tasks, while not requiring restructuring of code
- Goal: share MCU efficiently
 - Embedded Systems: To simplify our program design by allowing us to partition design into multiple independent components
 - PCs/Workstations/Servers: To allow multiple users to share a computer system

Example: Secure Answering Machine (SAM)³



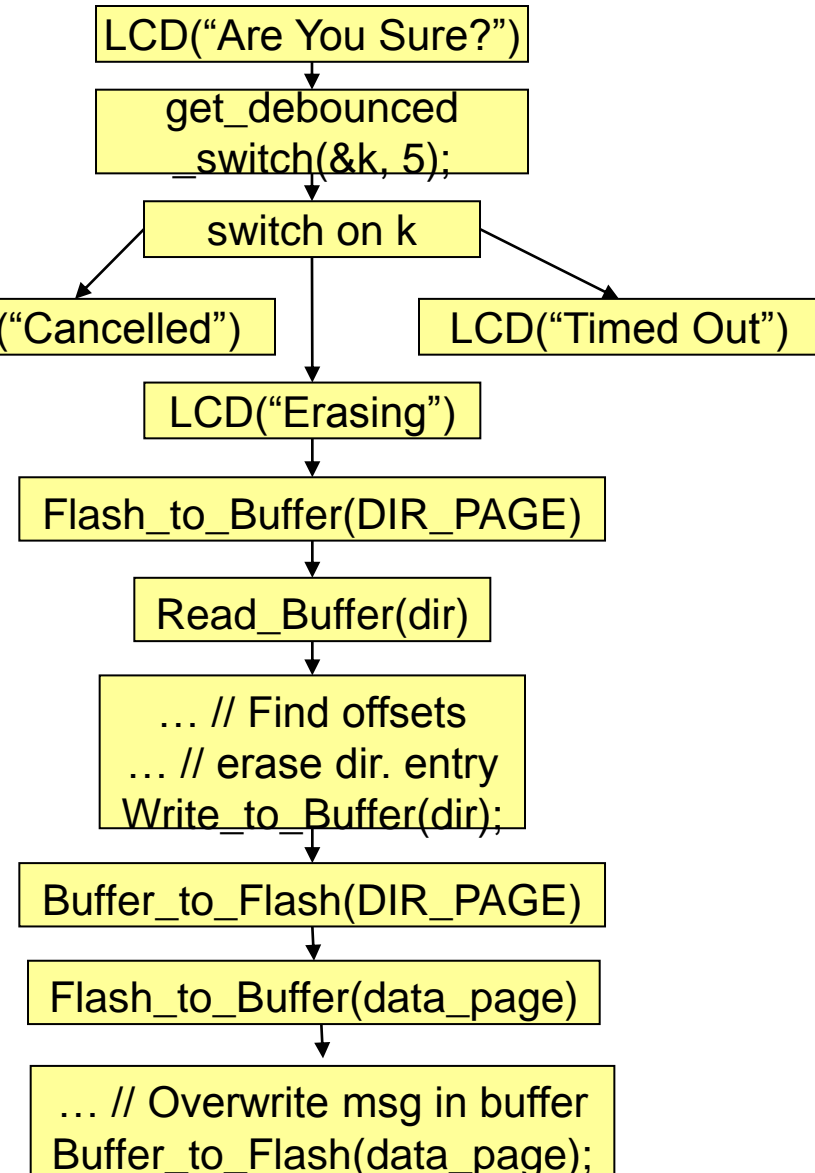
- *Testing the limits of our non-preemptive run-to-completion scheduler*
- Secure Answering Machine
 - Stores encrypted voice messages in serial Flash memory
 - Want to delete messages fully, not just remove entry from directory (as with file systems for PCs)
 - Also have a user interface: LCD, switches

SAM Delete Function and Timing

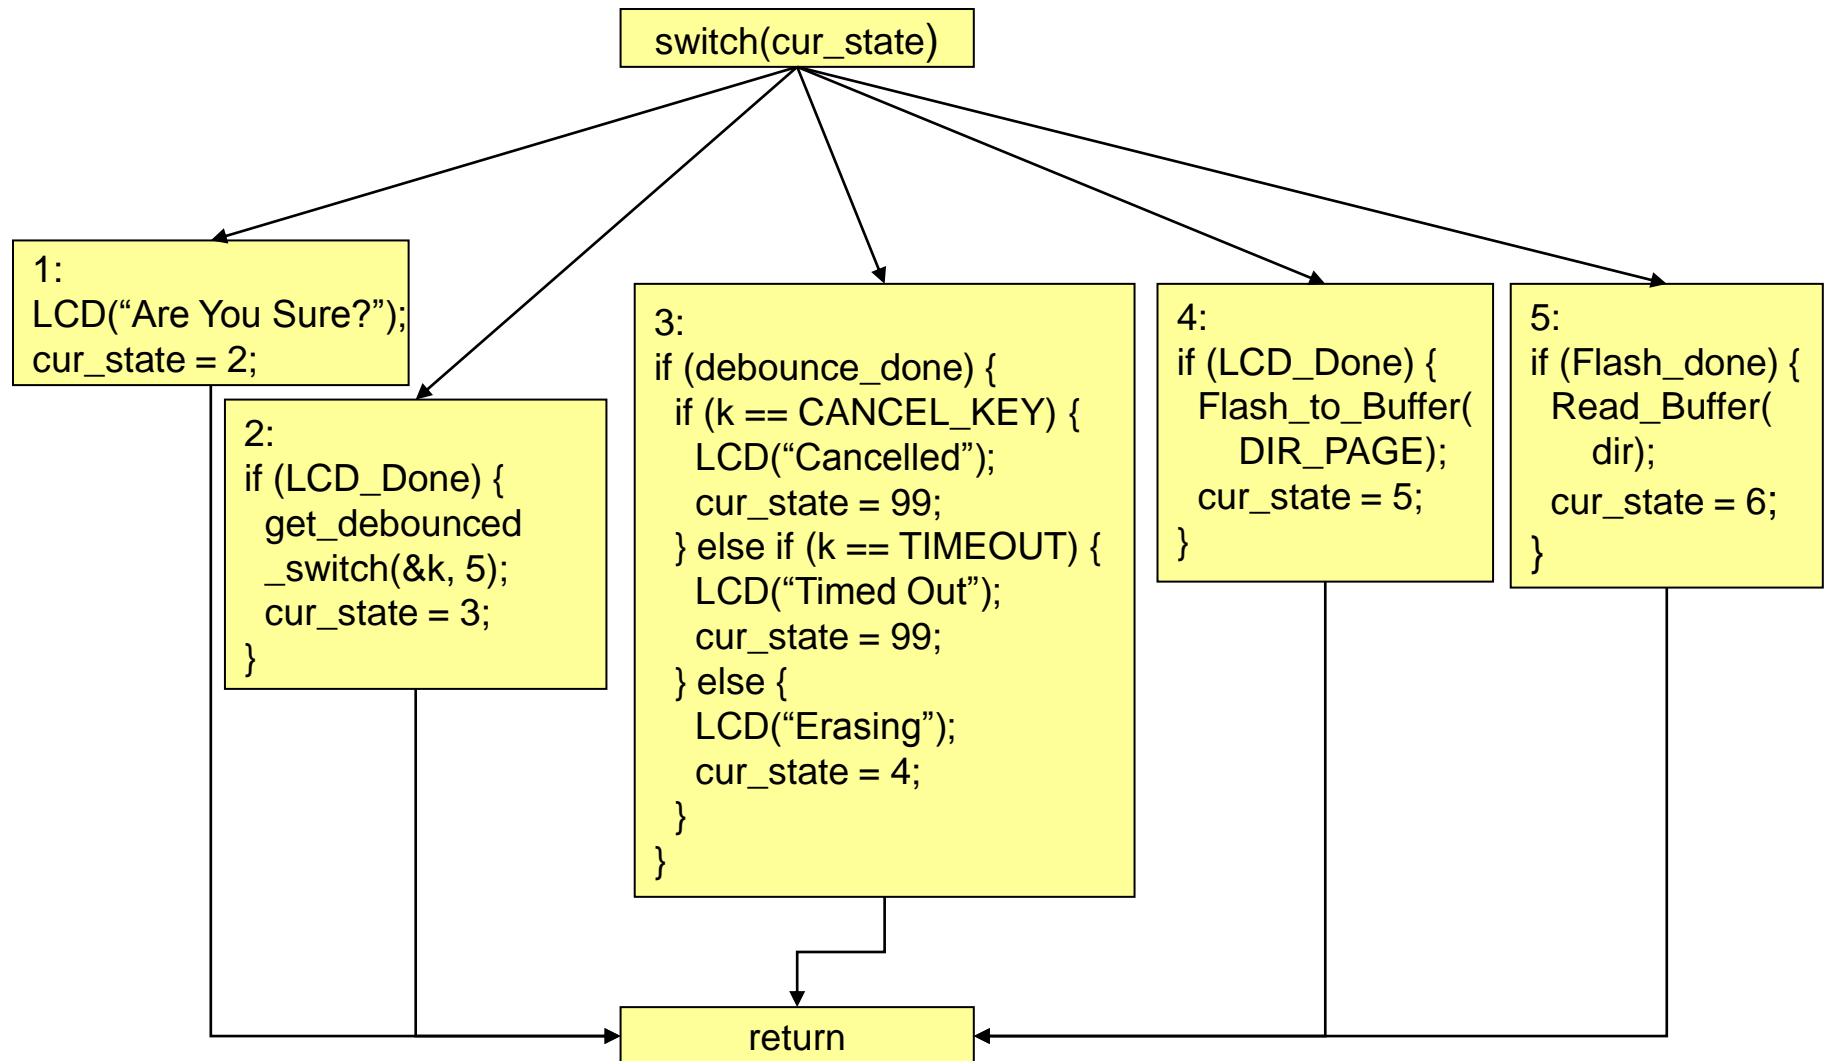
```
void Delete_Message(unsigned mes_num) {  
    ...  
    LCD("Are you sure?");           // 10 ms  
    get_debounced_switch(&k, 5);    // 400 ms min, 5 s max  
    if (k == CANCEL_KEY) {  
        LCD("Cancelled");           // 10 ms  
    } else if (k == TIMEOUT) {  
        LCD("Timed Out");           // 10 ms  
    } else {  
        LCD("Erasing");             // 10 ms  
        Flash_to_Buffer(DIR_PAGE);  // 250 us  
        Read_Buffer(dir);           // 100 us  
        ...                         // find offsets  
        ...                         // erase dir. entry  
        Write_to_Buffer(dir);       // 6 us  
        Buffer_to_Flash(DIR_PAGE);   // 20 ms  
        Flash_to_Buffer(data_page);  
        ...                         // overwrite msg: 50 us  
        Buffer_to_Flash(data_page);  // 20 ms  
        LCD("Done");  
    }  
}
```

How to do with the RTC Scheduler?

- Since task must **Run To Completion...**
- The delete function could take up to five seconds to run, halting all other tasks (but interrupts run)
- Other software needs to keep running, so break this into pieces. Run one piece at a time.
- How to split?
 - Each piece ends where processor waits for user (e.g. debounced switch) or other devices (Flash, LCD).
- How to control execution of pieces?
 1. Use a task per piece, use calls to `Reschedule_Task` and `Disable_Task` as needed
 - Need 13 different tasks (12 shown here)
 2. Use a state machine within one task



State Machine in One Task



Daydreaming

- Some functions are causing trouble for us – they use slow devices which make the processor wait
 - LCD: controller chip on LCD is slow
 - DataFlash: it takes time to program Flash EEPROM
 - Switch debouncing: physical characteristics of switch, time-outs
- Wouldn't it be great if we could ...
 - Make those slow functions *yield the processor* to other tasks?
 - Not have the processor start running that code again *until the device is ready*?
 - Maybe even have the processor interrupt less-important tasks?
 - *Avoid breaking up one task* into many tasks, or a state machine?
 - Open ourselves up to a whole new species of bugs – *bugs which are very hard to duplicate and track down*?

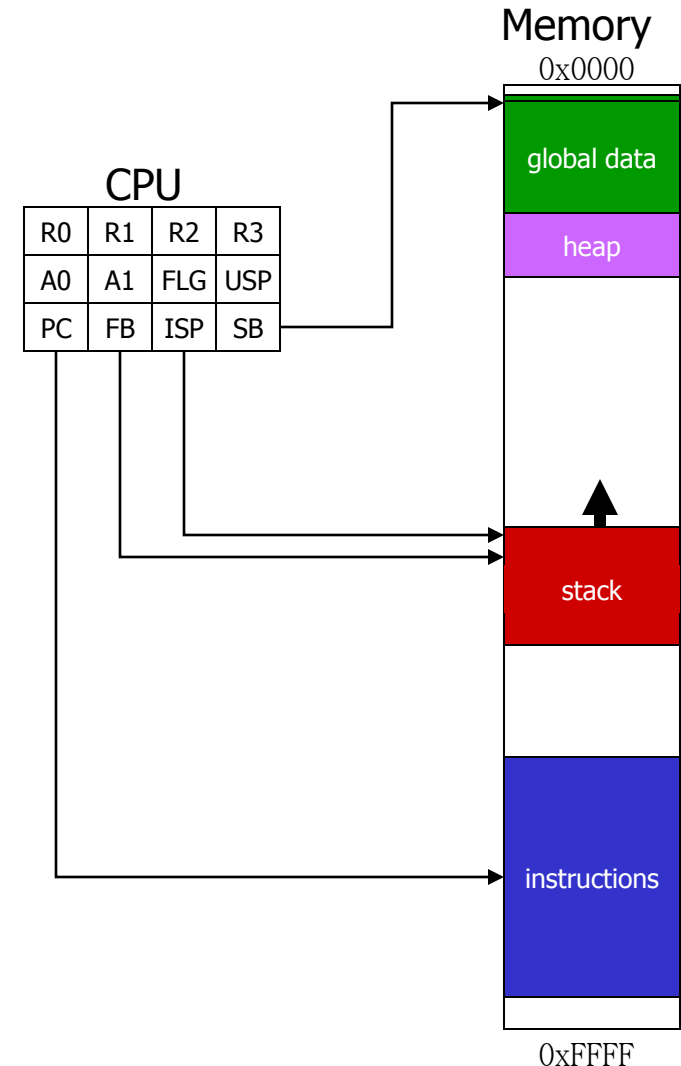


Preemptive Scheduling Kernel

- What we need is a *kernel*
 - Shares the processor among multiple concurrently running tasks/threads/processes
 - Can forcibly switch the processor from thread A to B and resume A later (preemption)
 - Can resume threads when their data is ready
 - Can simplify inter-thread communication by providing mechanisms
 - The heart of any operating system
- Terminology: “Kernel Mode”
 - PCs and workstations don’t expose all of the machine to the user’s program
 - Only code in *kernel* or *supervisor* mode have full access
 - Some high-end embedded processors have a restricted mode (e.g. ARM, MIPS)

What Execution State Information Exists?

- A program, process or thread in execution which has *state information*...
 - Current instruction – identified with program counter
 - Call stack – identified with stack pointer
 - Arguments, local variables, return addresses, dynamic links
 - Other CPU state
 - Register values (anything which will be shared and could be affected by the other processes) – general purpose registers, stack pointer, etc.
 - Status flags (zero, carry, interrupts enabled, carry bit, etc.)
 - Other information as well
 - Open files, memory management info, process number, scheduling information
 - Ignore for now

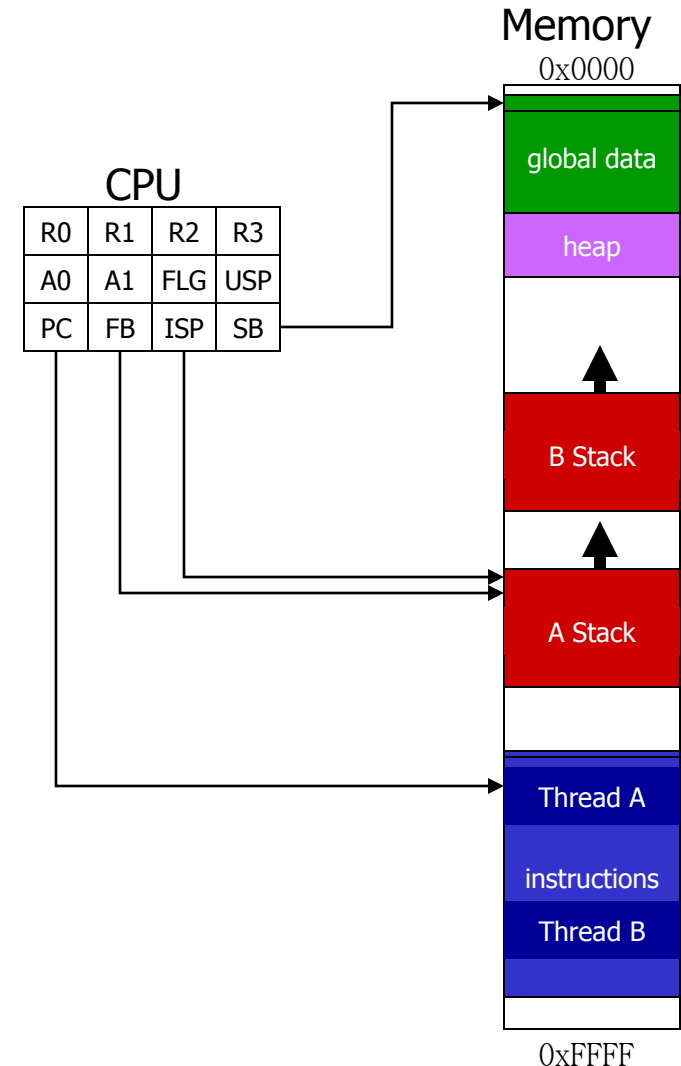


Processes vs. Threads

- Process – No information is visible to other processes (nothing is shared)
- Thread – Shares address space and code with other threads (also called *lightweight process*)
- One big side effect: context switching time varies
 - Switching among processes requires swapping large amounts of information
 - Switching among threads requires swapping much less information (PC, stack pointer and other registers, CPU state) and is much faster
- For this discussion, concepts apply equally to threads and processes

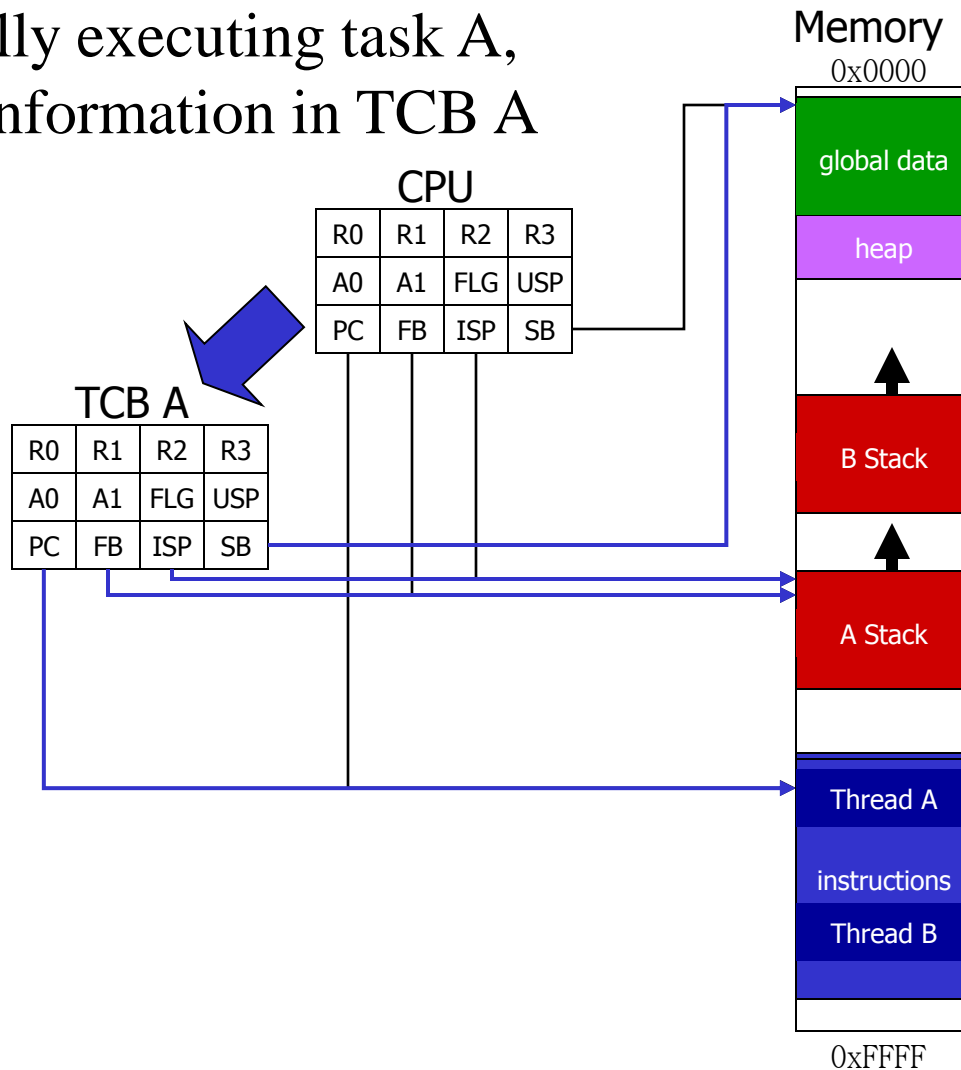
Maintaining State for Multiple Threads

- Store this thread-related information in a task/thread control block (TCB)
 - process control block = PCB
- Shuffling information between CPU and multiple TCBs lets us share processor
- Consider case of switching from thread A to thread B
 - Assume we have a call stack for each thread
 - Assume we can share global variables among the two threads
 - Standard for threads
 - For MSP430 architecture, SB register is same for both threads



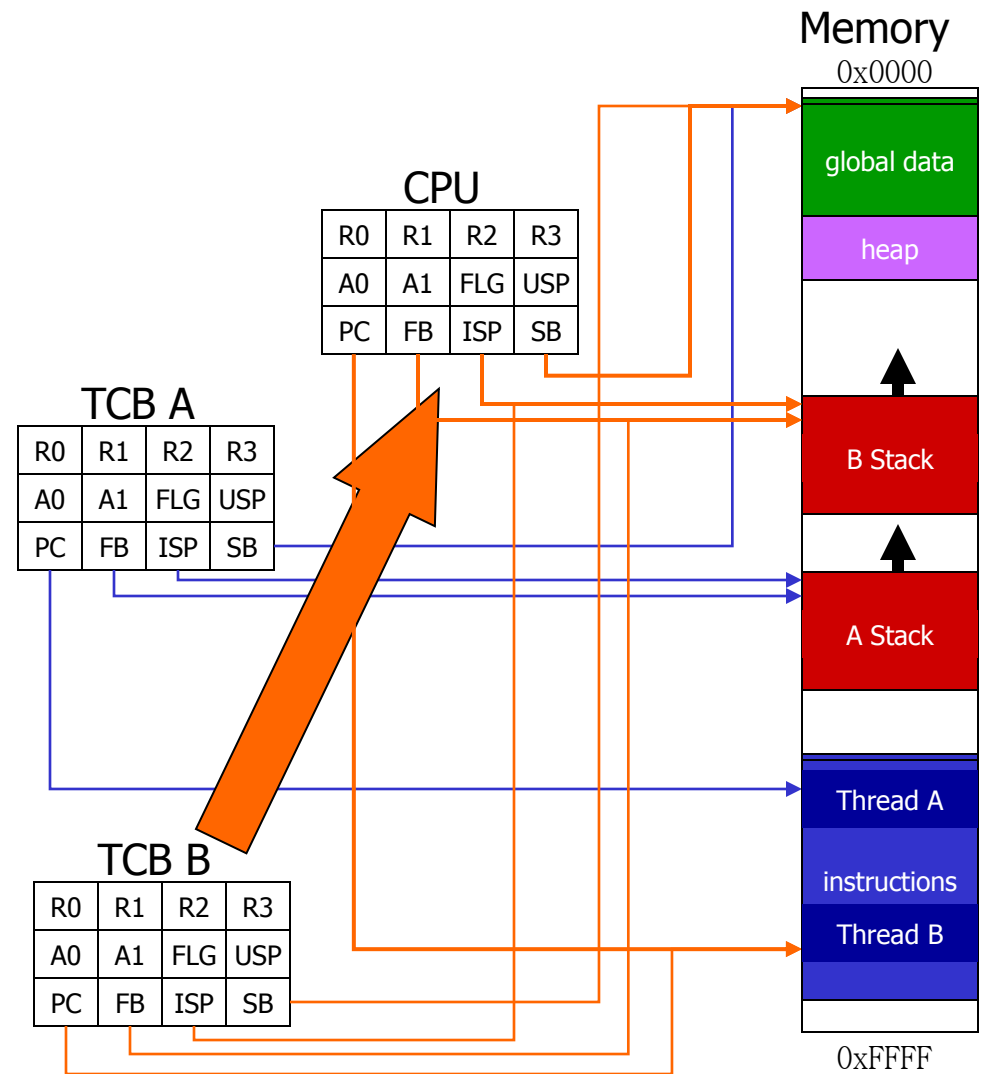
Step 1. Copy CPU State into TCB A

CPU is initially executing task A,
so save this information in TCB A



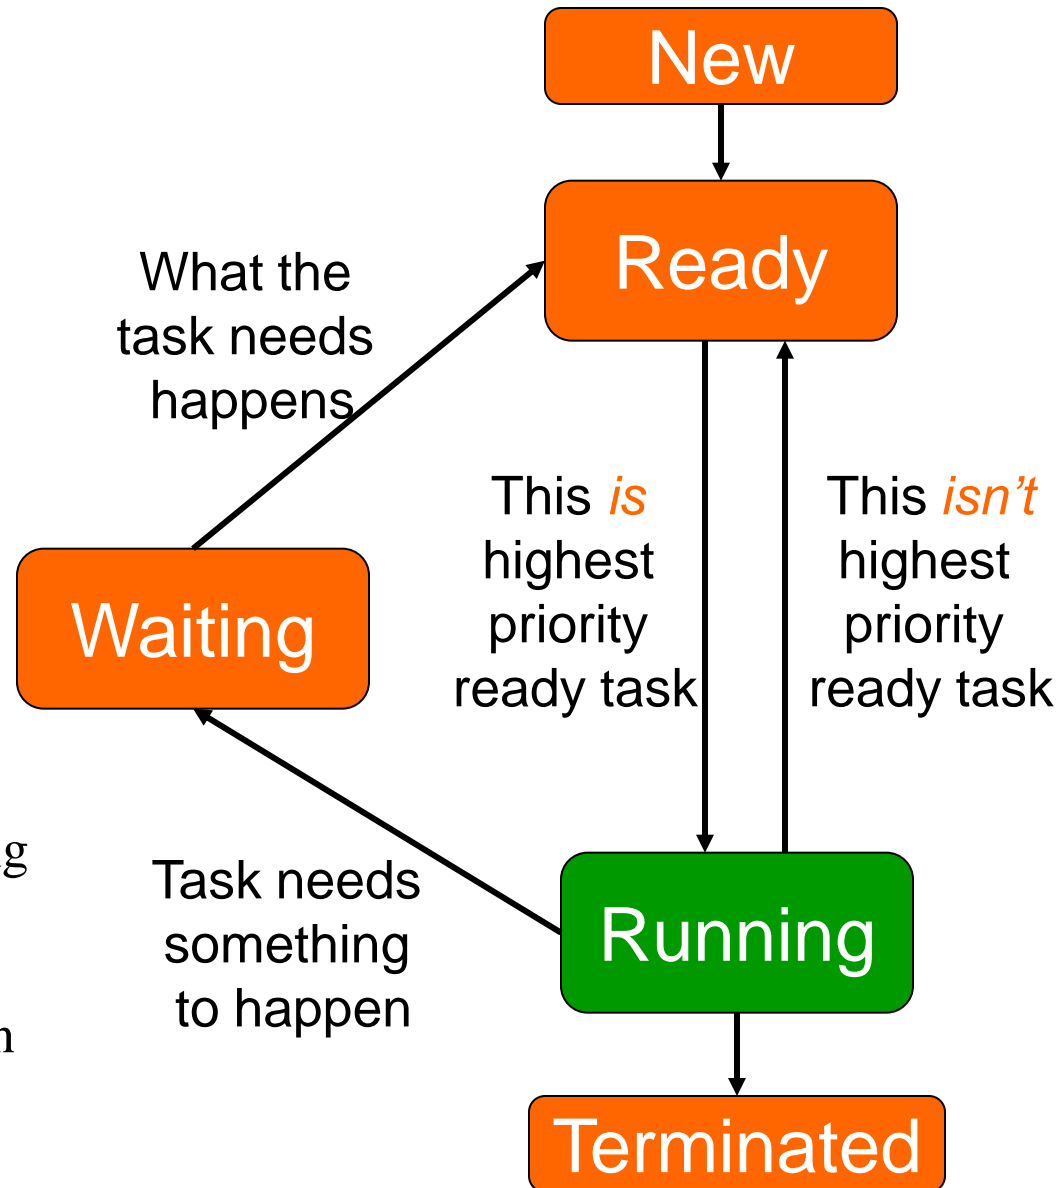
Step 2. Reload Old CPU State from TCB B

- Reloading a previously saved state configures the CPU to execute task B from where it left off
- This *context switching* is performed by the *dispatcher* code
- Dispatcher is typically written in assembly language to gain access to registers not visible to C programmer



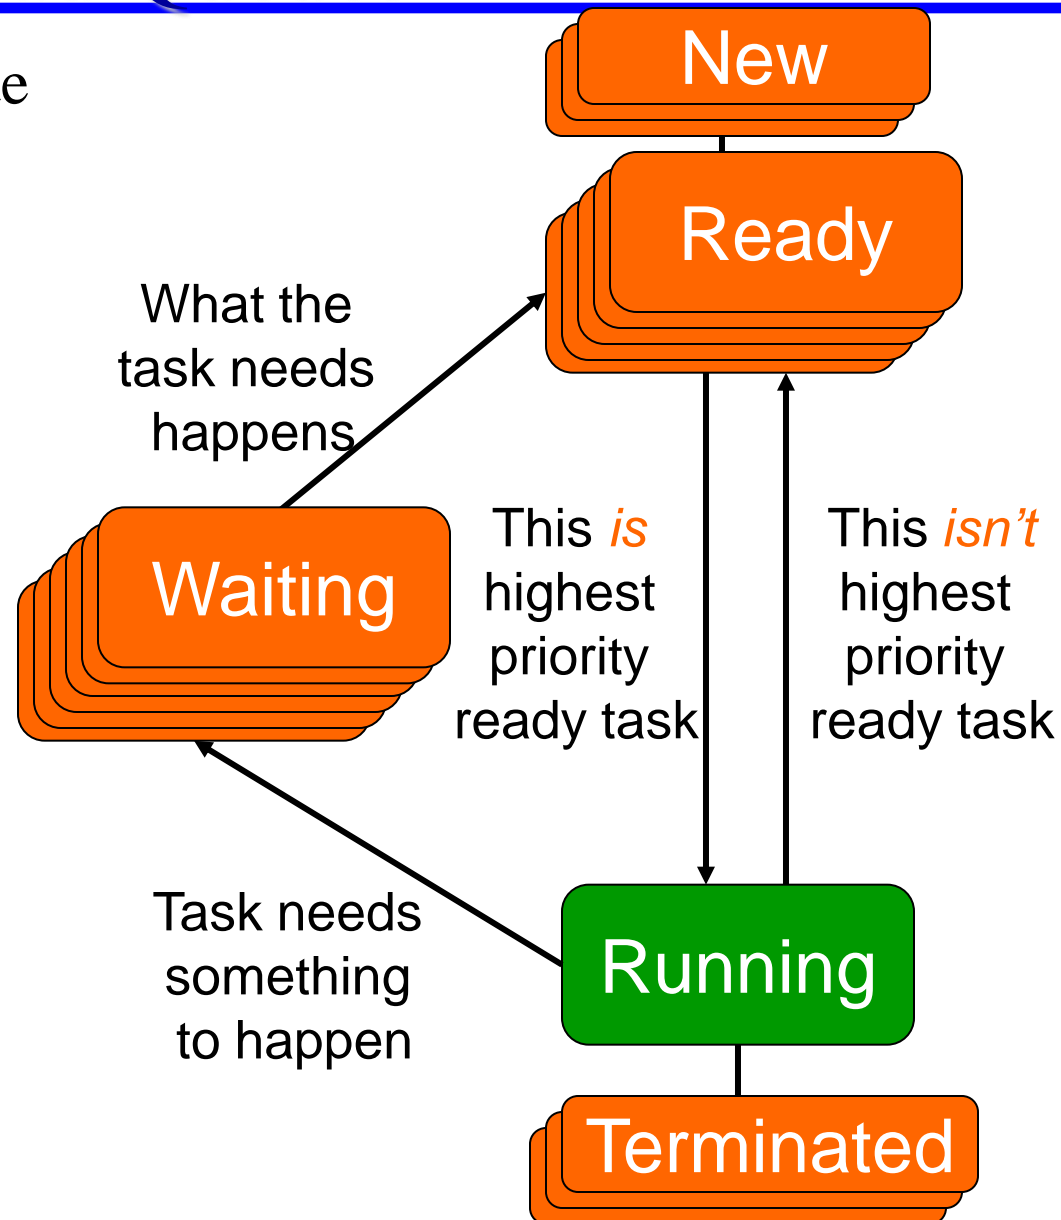
Thread States

- Now that we can share the CPU, let's do it!
- Define five possible states for a thread to be in
 - New – just created, but not running yet
 - Running – instructions are being executed (only one thread can be running at a time!)
 - Waiting/Blocking – thread is waiting for an event to occur
 - Ready – process is not waiting but not running yet (is a candidate for running)
 - Terminated – process will run no more



Thread Queues

- Create a queue for each state (except running)
- Now we can store thread control blocks in the appropriate queues
- Kernel moves tasks among queues/processor registers as needed



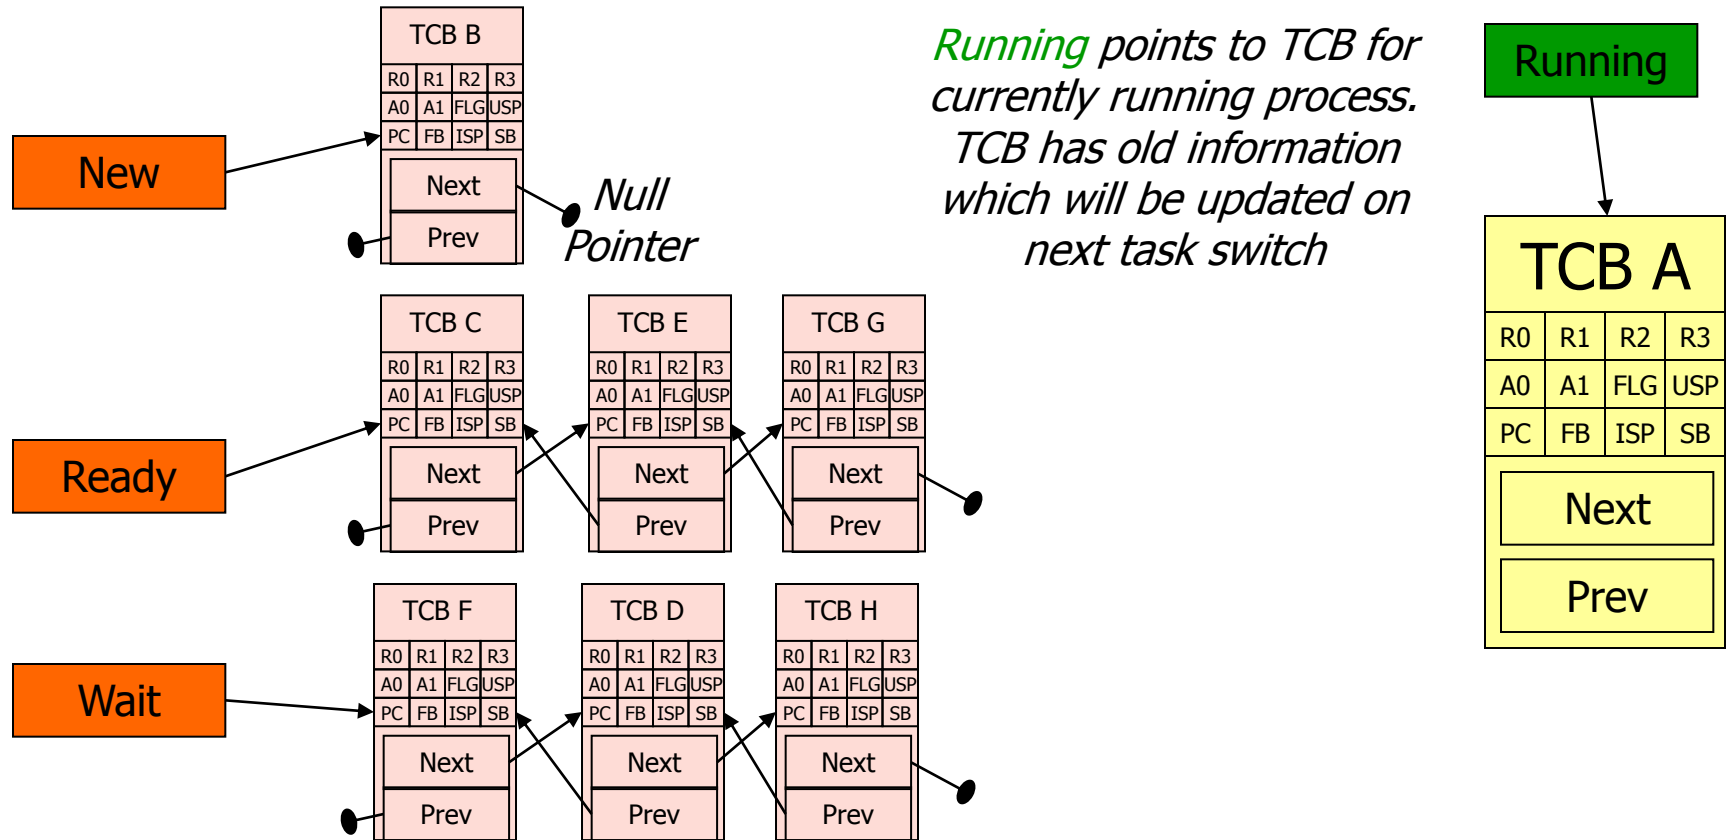
Preemptive vs. Non-Preemptive

- Non-preemptive kernel/cooperative multitasking
 - Each task must explicitly give up control of CPU
 - E.g. return from task code, call yield function
 - Asynchronous events are handled by ISRs
 - ISR always returns to interrupted task
 - Can use non-reentrant code (covered later)
 - Task level response time can be slower as slowest task must complete
 - Generally don't need semaphores
- Preemptive kernel
 - At each scheduling point, the highest priority task ready to run is given CPU control
 - If a higher priority task becomes ready, the currently running task is suspended and moved to the ready queue
 - Maximum response time is less than in non-preemptive system
 - Non-reentrant code should not be used
 - Shared data typically needs semaphores

Thread State Control

- Use OS scheduler to keep track of threads and their states
 - For each state, OS keeps a queue of TCBs for all processes in that state
 - Moves TCBs between queues as thread state changes
 - OS's scheduler chooses among *Ready* threads for execution based on priority
 - Scheduling Rules
 - Only the thread itself can decide it should be *waiting (blocked)*
 - A *waiting* thread never gets the CPU. It must be signaled by an ISR or another thread.
 - Only the scheduler moves tasks between *ready* and *running*
- What changes the state of a thread?
 - The OS receives a timer tick which forces it to decide what to run next
 - The thread voluntarily yields control
 - The thread requests information which isn't ready yet

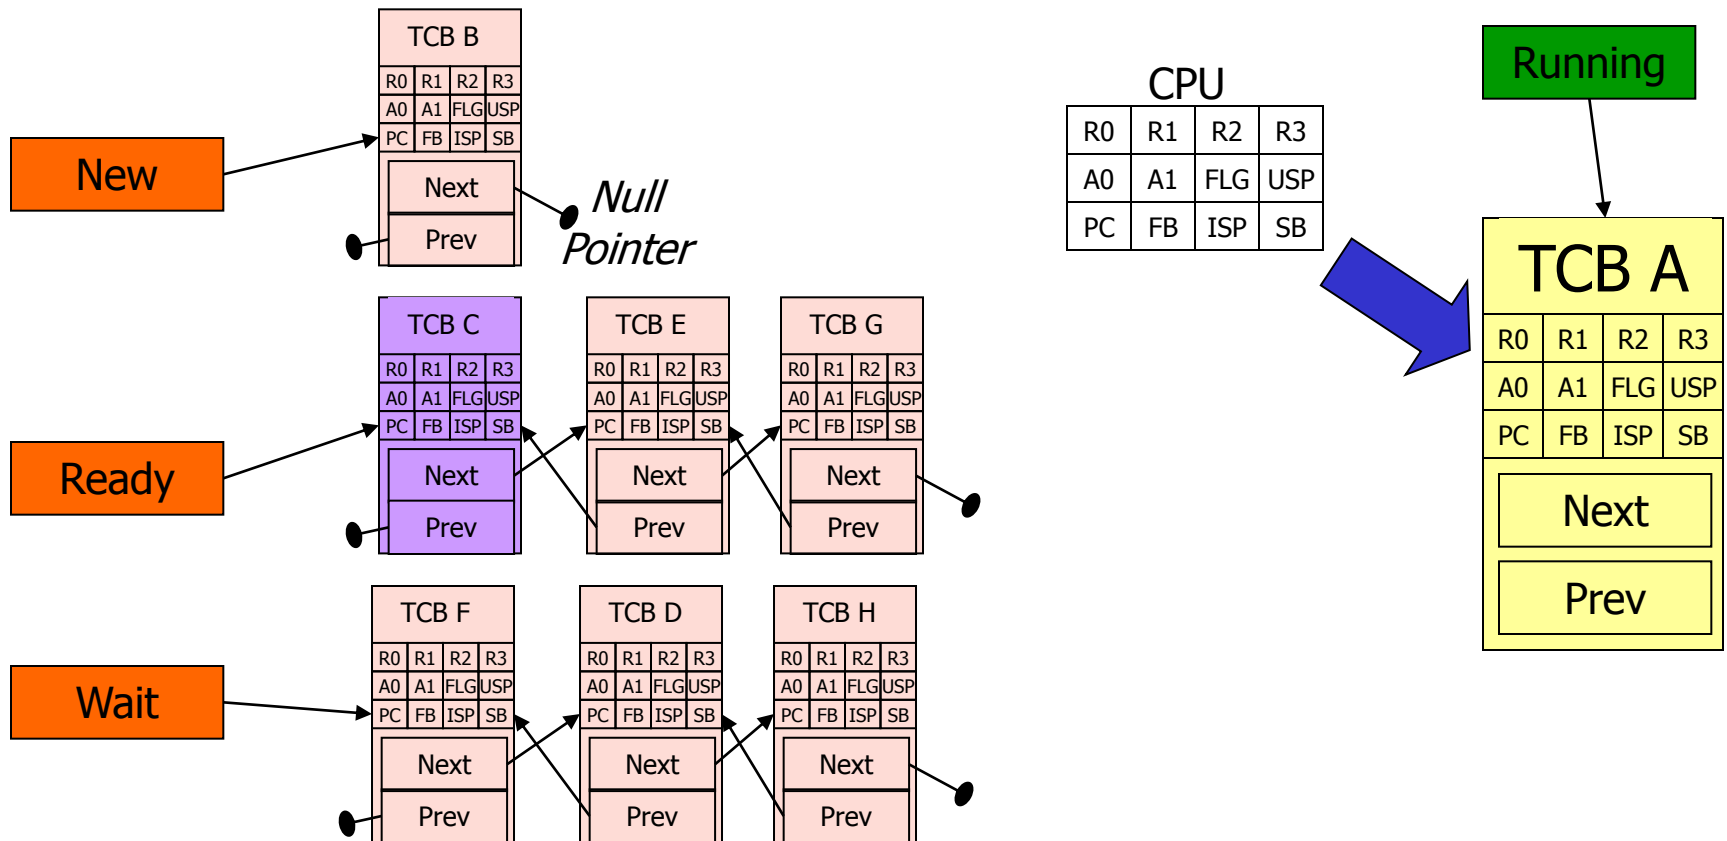
Overview of Data Structures for Scheduler



- Add Next, Prev pointers in each TCB to make it part of a doubly linked list
- Keep track of all TCBs
 - Create a pointer for each queue: Ready, Wait, New
 - Create a pointer for the currently running task's TCB

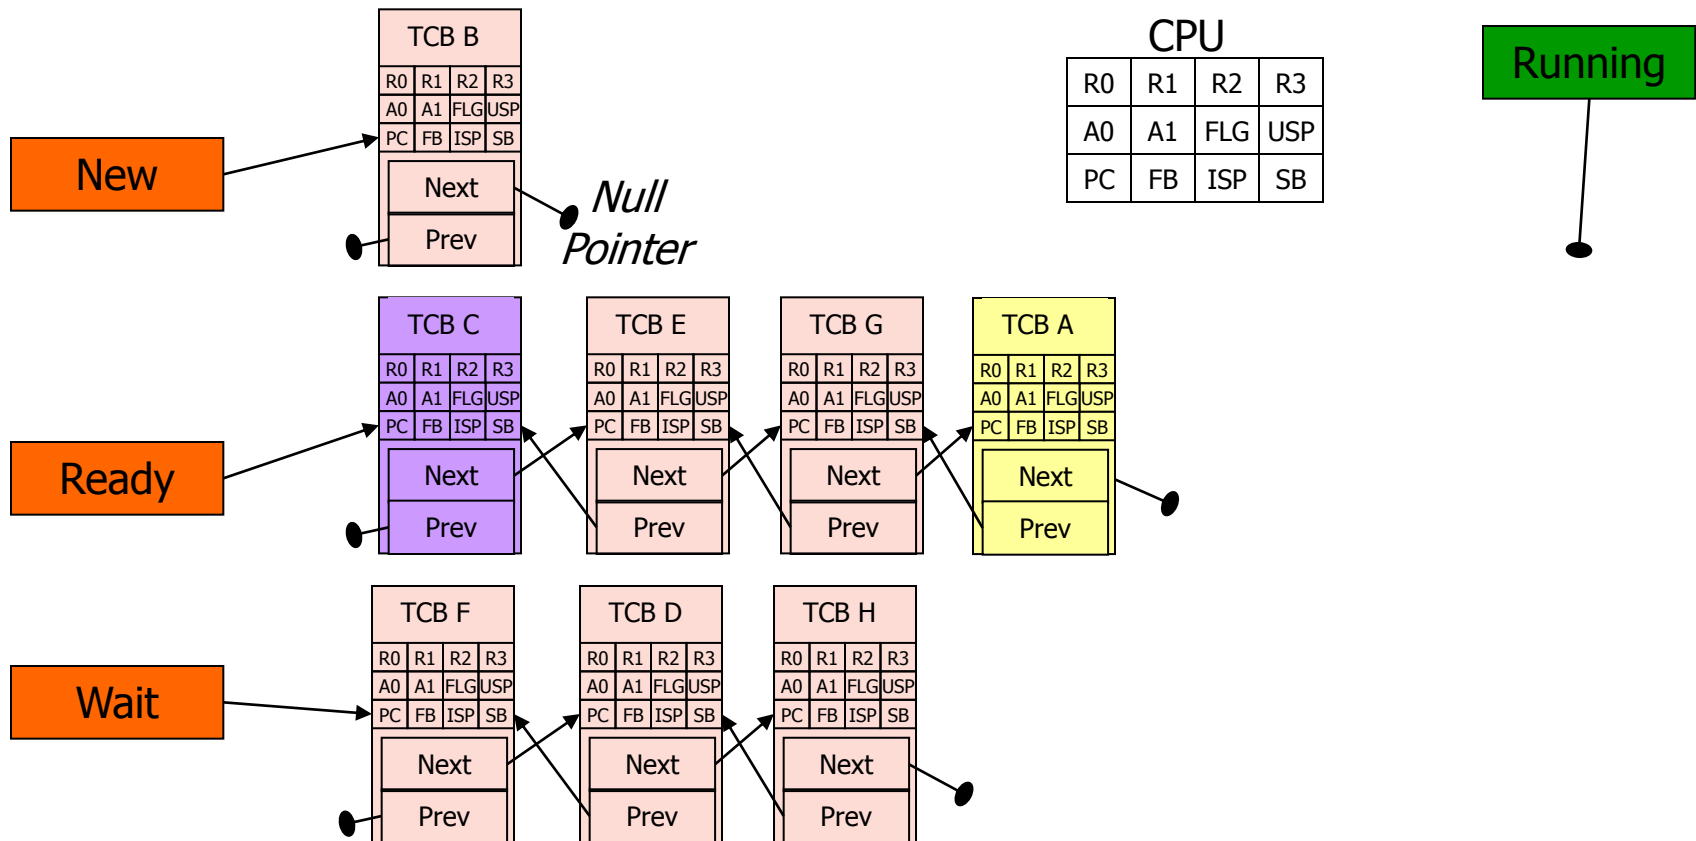
Example: Context Switch

- Thread A is running, and scheduler decides to run thread C instead. For example, thread A is still able to run, but has lower priority than thread C.
- Start by copying CPU state into TCB A



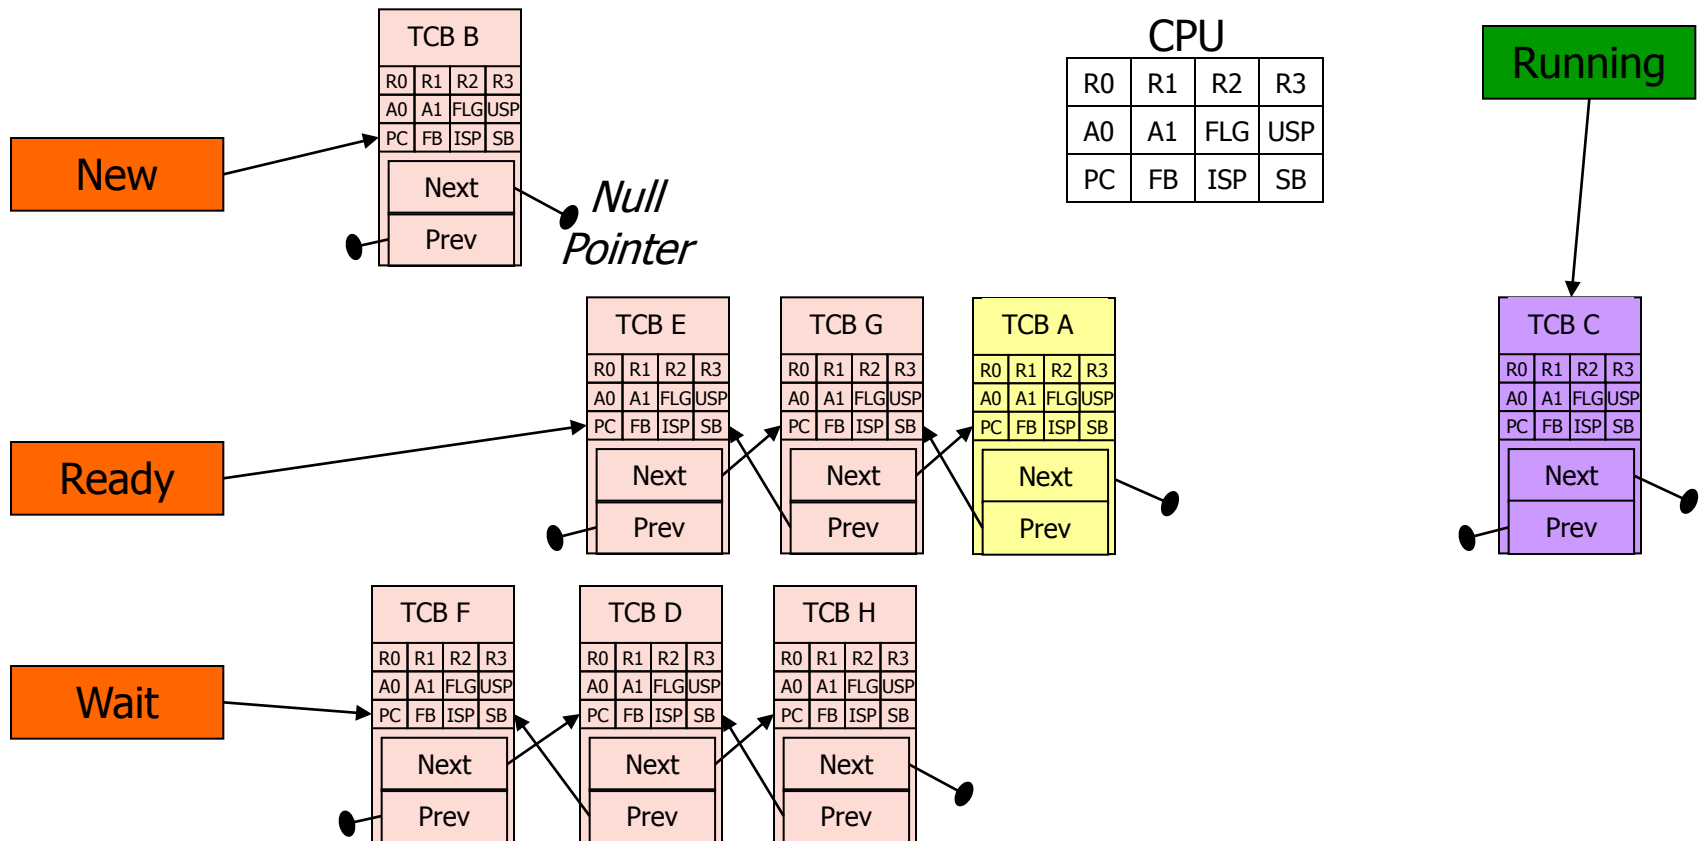
Example: Context Switch

- Insert TCB A into ready queue by modifying appropriate pointers



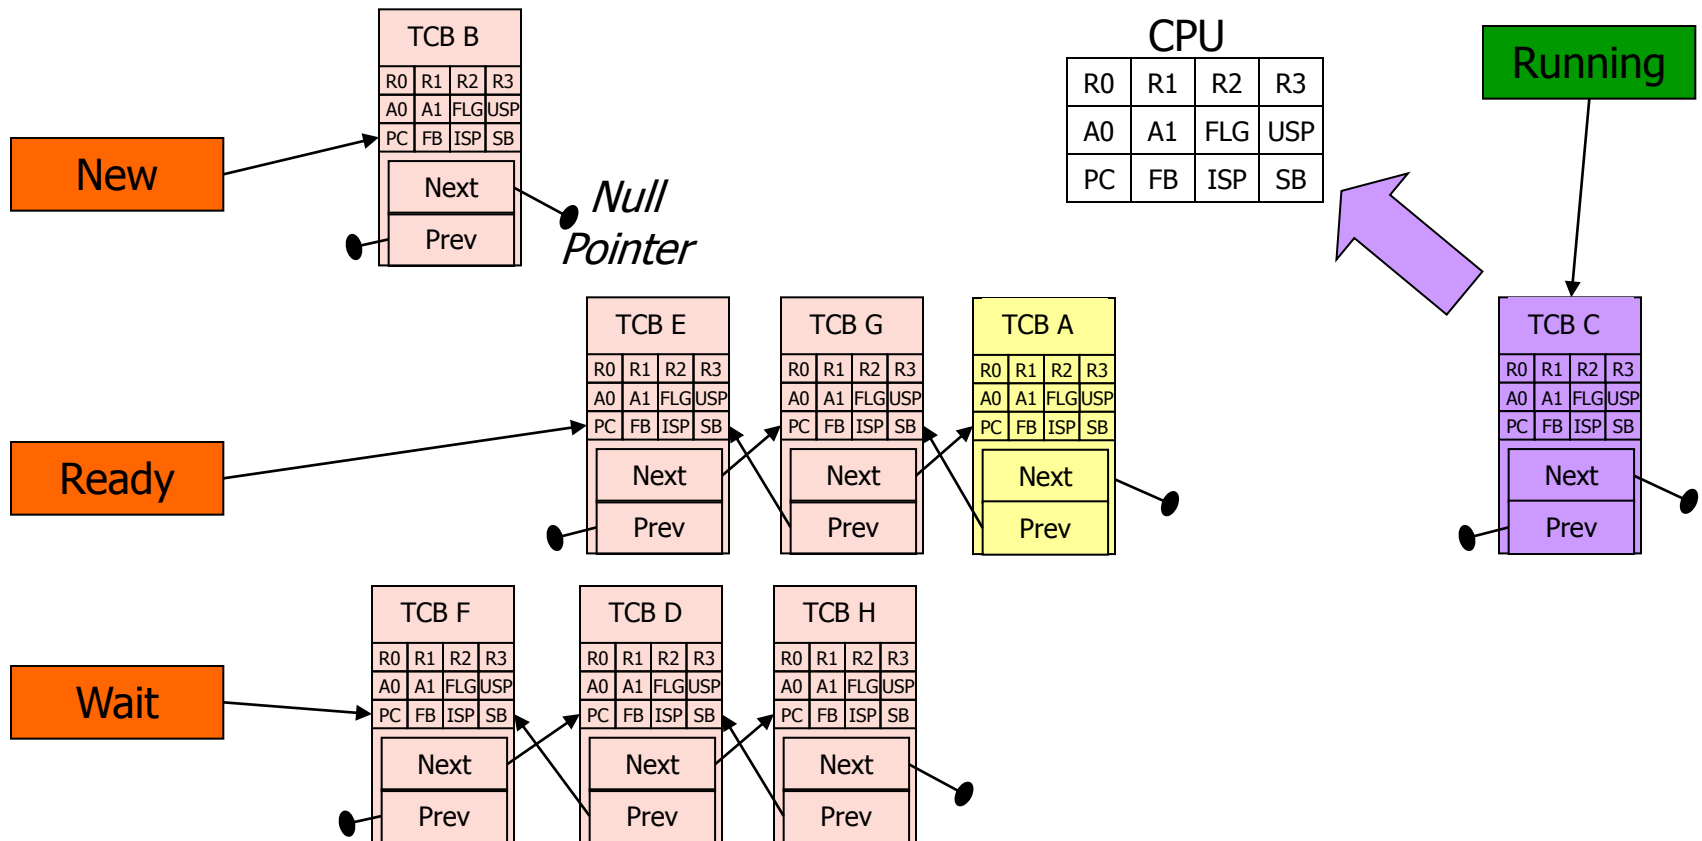
Example: Context Switch

- Remove thread C from the ready queue and mark it as the thread to run next



Example: Context Switch

- Copy thread C's state information back into the CPU and resume execution

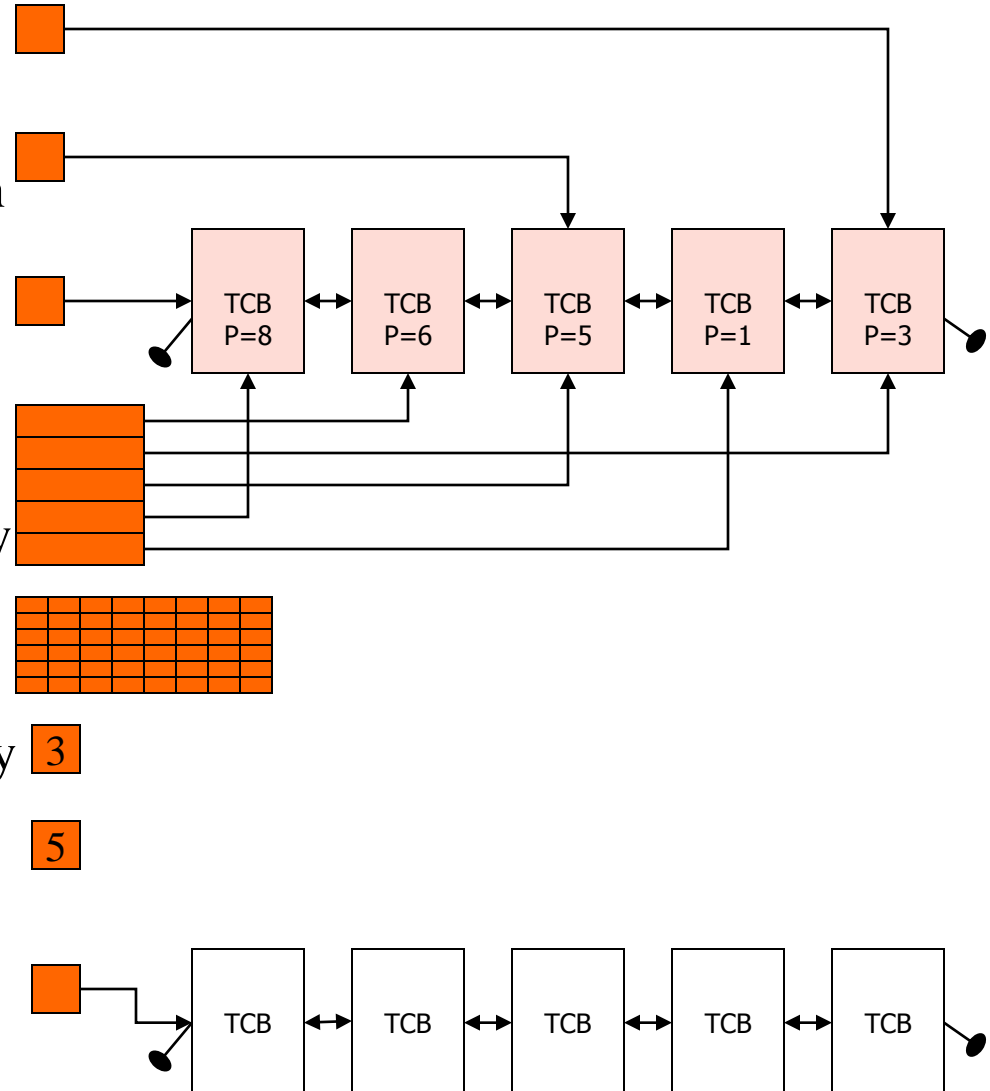


μC/OS-II

- Real-time kernel
 - Portable, scalable, preemptive RTOS
 - Ported to over 90 processors
- Pronounced “microC OS two”
- Written by Jean J. Labrosse of Micrium,
<http://ucos-ii.com>
- Implementation is different from material just presented for performance and feature reasons
 - CPU state is stored on thread’s own stack, not TCB
 - TCB keeps track of boundaries of stack space
 - TCB also tracks events and messages and time delays

Data Structures for $\mu\text{C}/\text{OS-II}$

- OSTCBCur - Pointer to TCB of currently running task
- OSTCBHighRdy - Pointer to highest priority TCB ready to run
- OSTCBLList - Pointer to doubly linked list of TCBs
- OSTCBPrioTbl[OS_LOWEST_PRIO + 1] - Table of pointers to created TCBs, ordered by priority
- OSReadyTbl - Encoded table of tasks ready to run
- OSPrioCur - Current task priority 3
- OSPrioHighRdy - Priority of highest ready task 5
- OSTCBFreeList - List of free OS_TCBs, use for creating new tasks



Dispatcher for μ C/OS-II

_OSctxSw:

PUSHM R0,R1,R2,R3,A0,A1,SB,FB

MOV.W _OSTCBCur, A0

;OSTCBCur->OSTCBStkPtr = Stack pointer

STC ISP, [A0]

;Call user definable OSTaskSwHook()

JSR _OSTaskSwHook

;OSTCBCur = OSTCBHighRdy

MOV.W _OSTCBHighRdy, _OSTCBCur

;OSPrioCur = OSPrioHighRdy

MOV.W _OSPrioHighRdy, _OSPrioCur

;Stack Pointer = OSTCBHighRdy->OSTCBStkPtr

MOV.W _OSTCBHighRdy, A0

LDC [A0], ISP

;Restore all processor registers from the new task's stack

POPM R0,R1,R2,R3,A0,A1,SB,FB

REIT