

Memory Expansion and Direct Memory Access

Lecture 22

In These Notes . . .

- Memory Types
- Memory Expansion
 - Interfacing
 - Parallel
 - Serial
- Direct Memory Access controllers

Memory Characteristics and Issues

- Volatility - Does it remember when the power fails?
- Persistence - How long does it remember *when there is power*?
- Speed – How quickly can it be read or written?
- Reprogrammability
 - Speed is often an issue for nonvolatile memories
 - Maximum number of W/E cycles
 - Programming voltage
- Cost
- Temperature Sensitivity - EPROMs forget at high temperatures

Types of Memory - ROM

- Mask
 - A custom ROM mask pattern is created
 - Large minimum order, NRE costs
- PROM
 - Program by burning fuses. Apply a high voltage for a certain amount of time.
 - Not erasable
- UV EPROM
 - Program by squeezing charge into floating transistor gate with high voltage
 - Erase entire memory at once with UV radiation

Types of Memory - EEPROM

- EEPROM
 - Can erase a byte at a time electrically
 - Limited reprogrammability: e.g. 100,000 cycles
 - Slow programming (up to 10 ms per byte)
- Flash EEPROM
 - Can erase entire chip, or just certain blocks
 - May have limited reprogrammability: e.g. 1,000 cycles for some low-end units
 - May have slow programming
 - Serial or parallel interface

Types of Memory - RAM

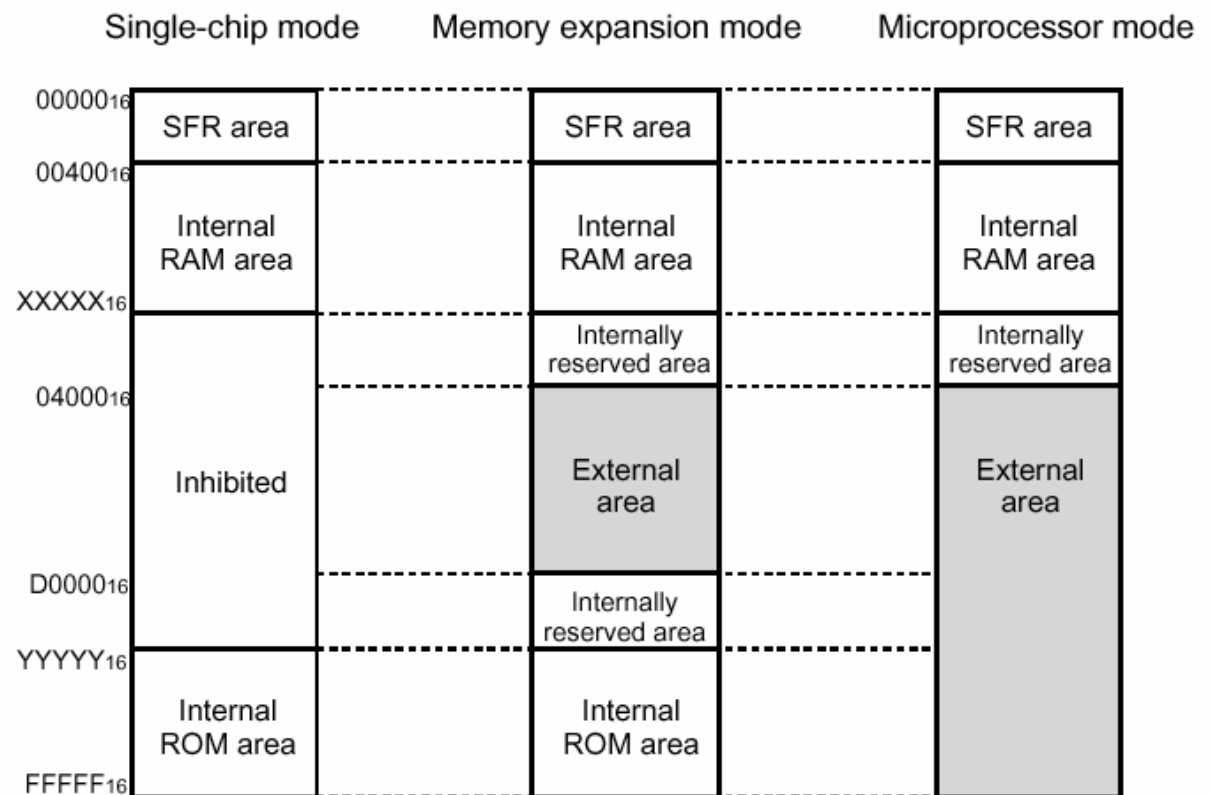
- SRAM
 - Use two or six transistors per bit
 - Fast - 10 ns or less
- DRAM
 - Use one transistor per bit
 - Acts like a capacitor, discharges in a few milliseconds
 - Incredibly cheap, but need to refresh each bit periodically
 - Slower than SRAM - 60 ns
 - Tricks to speed up access (e.g. page mode)

Parallel Memory Interface

- 1920 bytes are not enough for our data logger!
- We want to store up to 32 kBytes of data
- Choose a microcontroller with more memory?

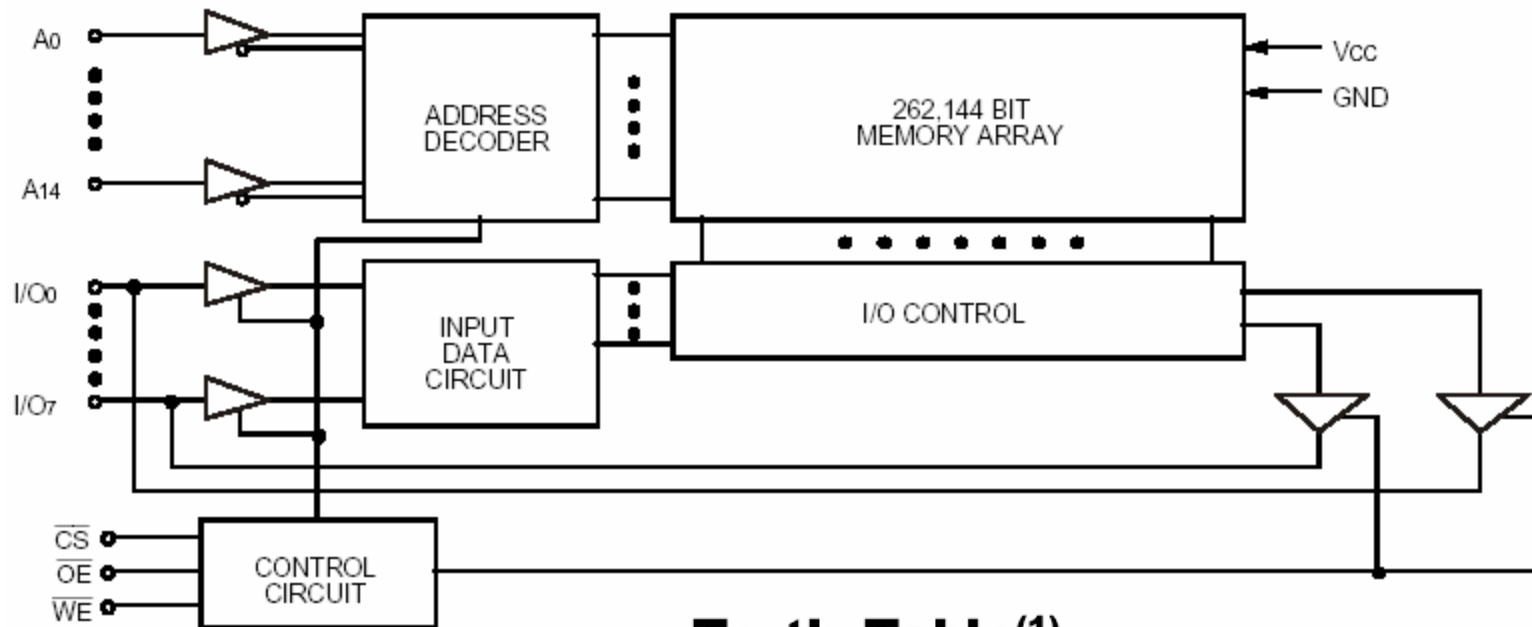
- Not available, or too expensive

- Instead choose a microcontroller with a memory expansion mode
 - M16C/62



Modes for M16C/62 MCU

SRAM Example – IDT71256L



- 32k x 8 SRAM

Truth Table⁽¹⁾

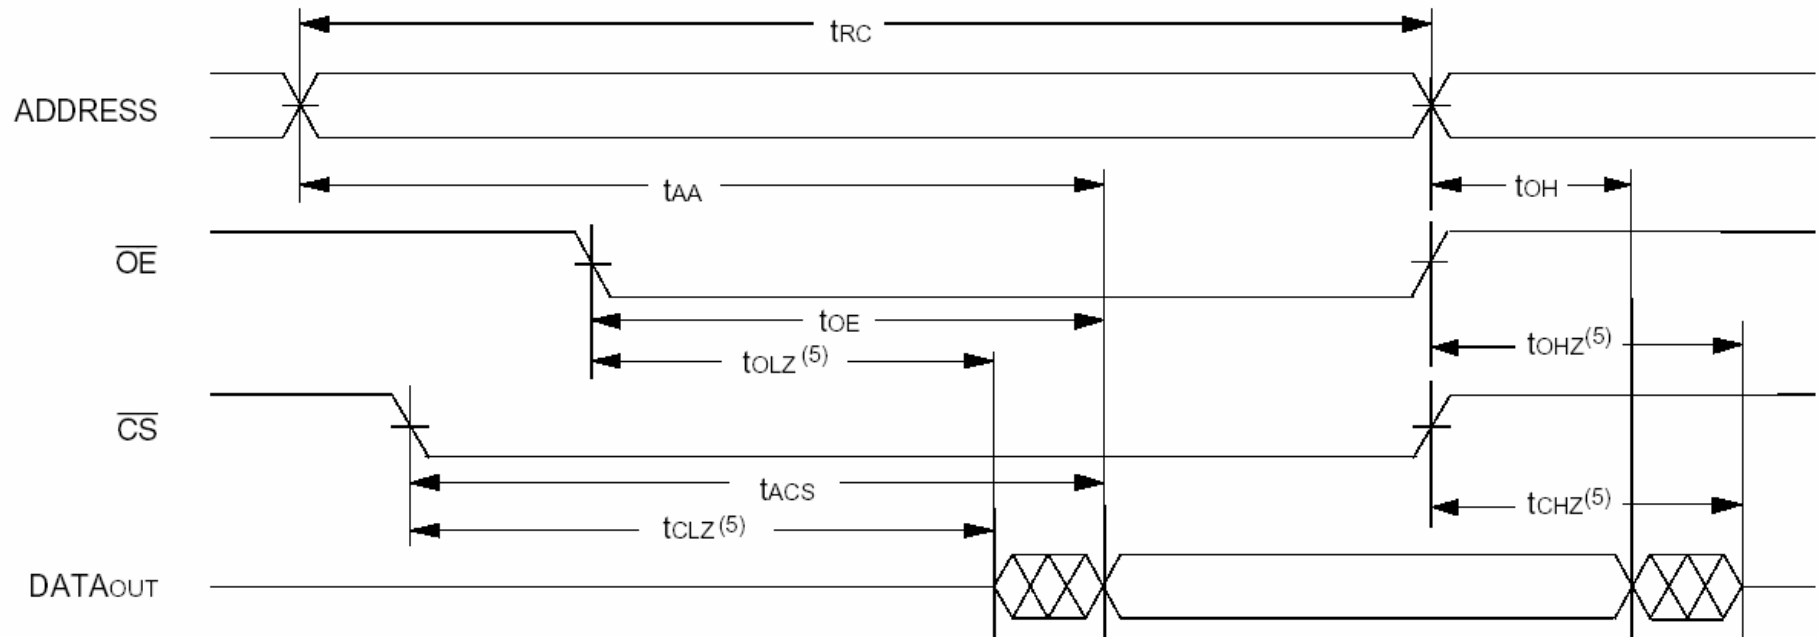
\overline{WE}	\overline{CS}	\overline{OE}	I/O	Function
X	H	X	High-Z	Standby (ISB)
X	V _{HC}	X	High-Z	Standby (ISB1)
H	L	H	High-Z	Output Disabled
H	L	L	DOUT	Read Data
L	L	X	DIN	Write Data

NOTE:

1. H = V_{IH}, L = V_{IL}, X = Don't care.

IDT71256L Read Timing

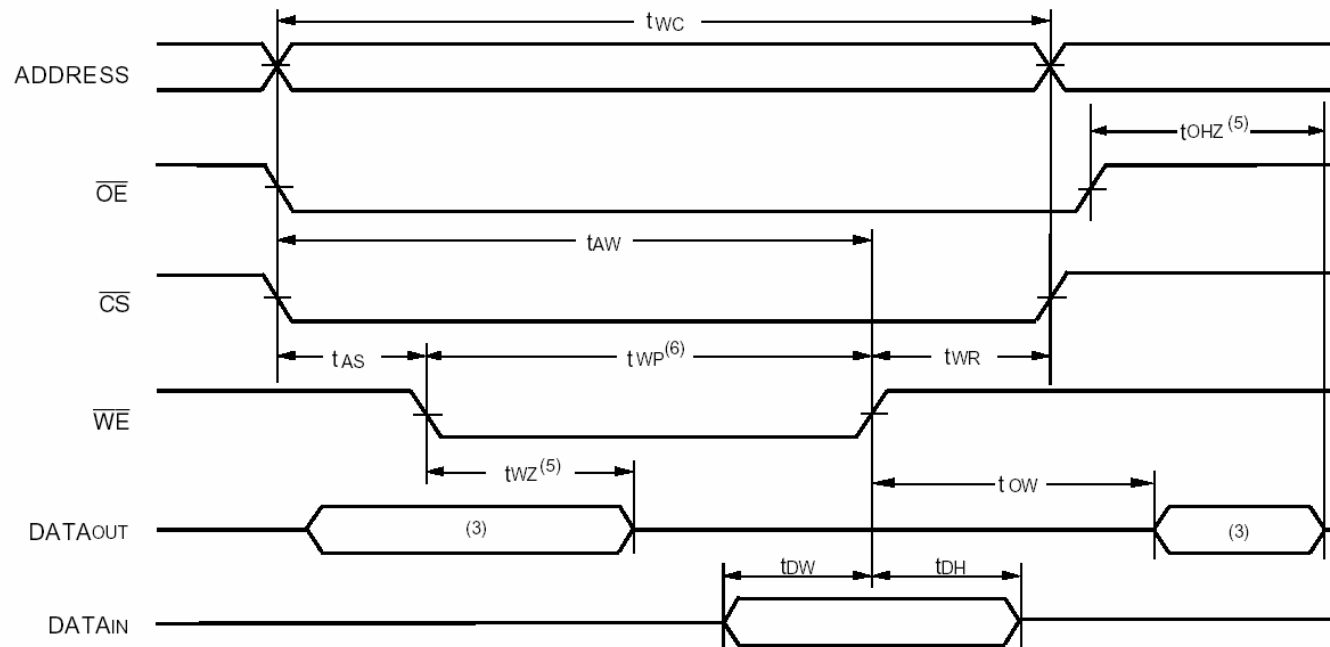
Timing Waveform of Read Cycle No. 1⁽¹⁾



Symbol	Parameter	Minimum (ns)	Maximum (ns)
t_{RC}	Read Cycle Time	20	-
t_{AA}	Address Access time	-	20
t_{OLZ}	Output Enable to Output in Low-Z	2	-
t_{OE}	Output Enable to Output Valid	-	10
t_{OHZ}	Output Disable to Output in Hi-Z	2	8


IDT71256L Write Timing

Timing Waveform of Write Cycle No. 1 ($\overline{\text{WE}}$ Controlled Timing)^(1,2,4,6)



Symbol	Parameter	Minimum (ns)	Maximum (ns)
t_{WC}	Write Cycle Time	20	-
t_{AW}	Address Valid to End-of-Write	15	-
t_{WP}	Write Pulse Width	15	-
t_{AS}	Address Set-Up Time	0	-
t_{DW}	Data to Write Time Overlap	11	-
t_{DH}	Data Hold from Write Time	0	-


External Memory Access – Separate Buses



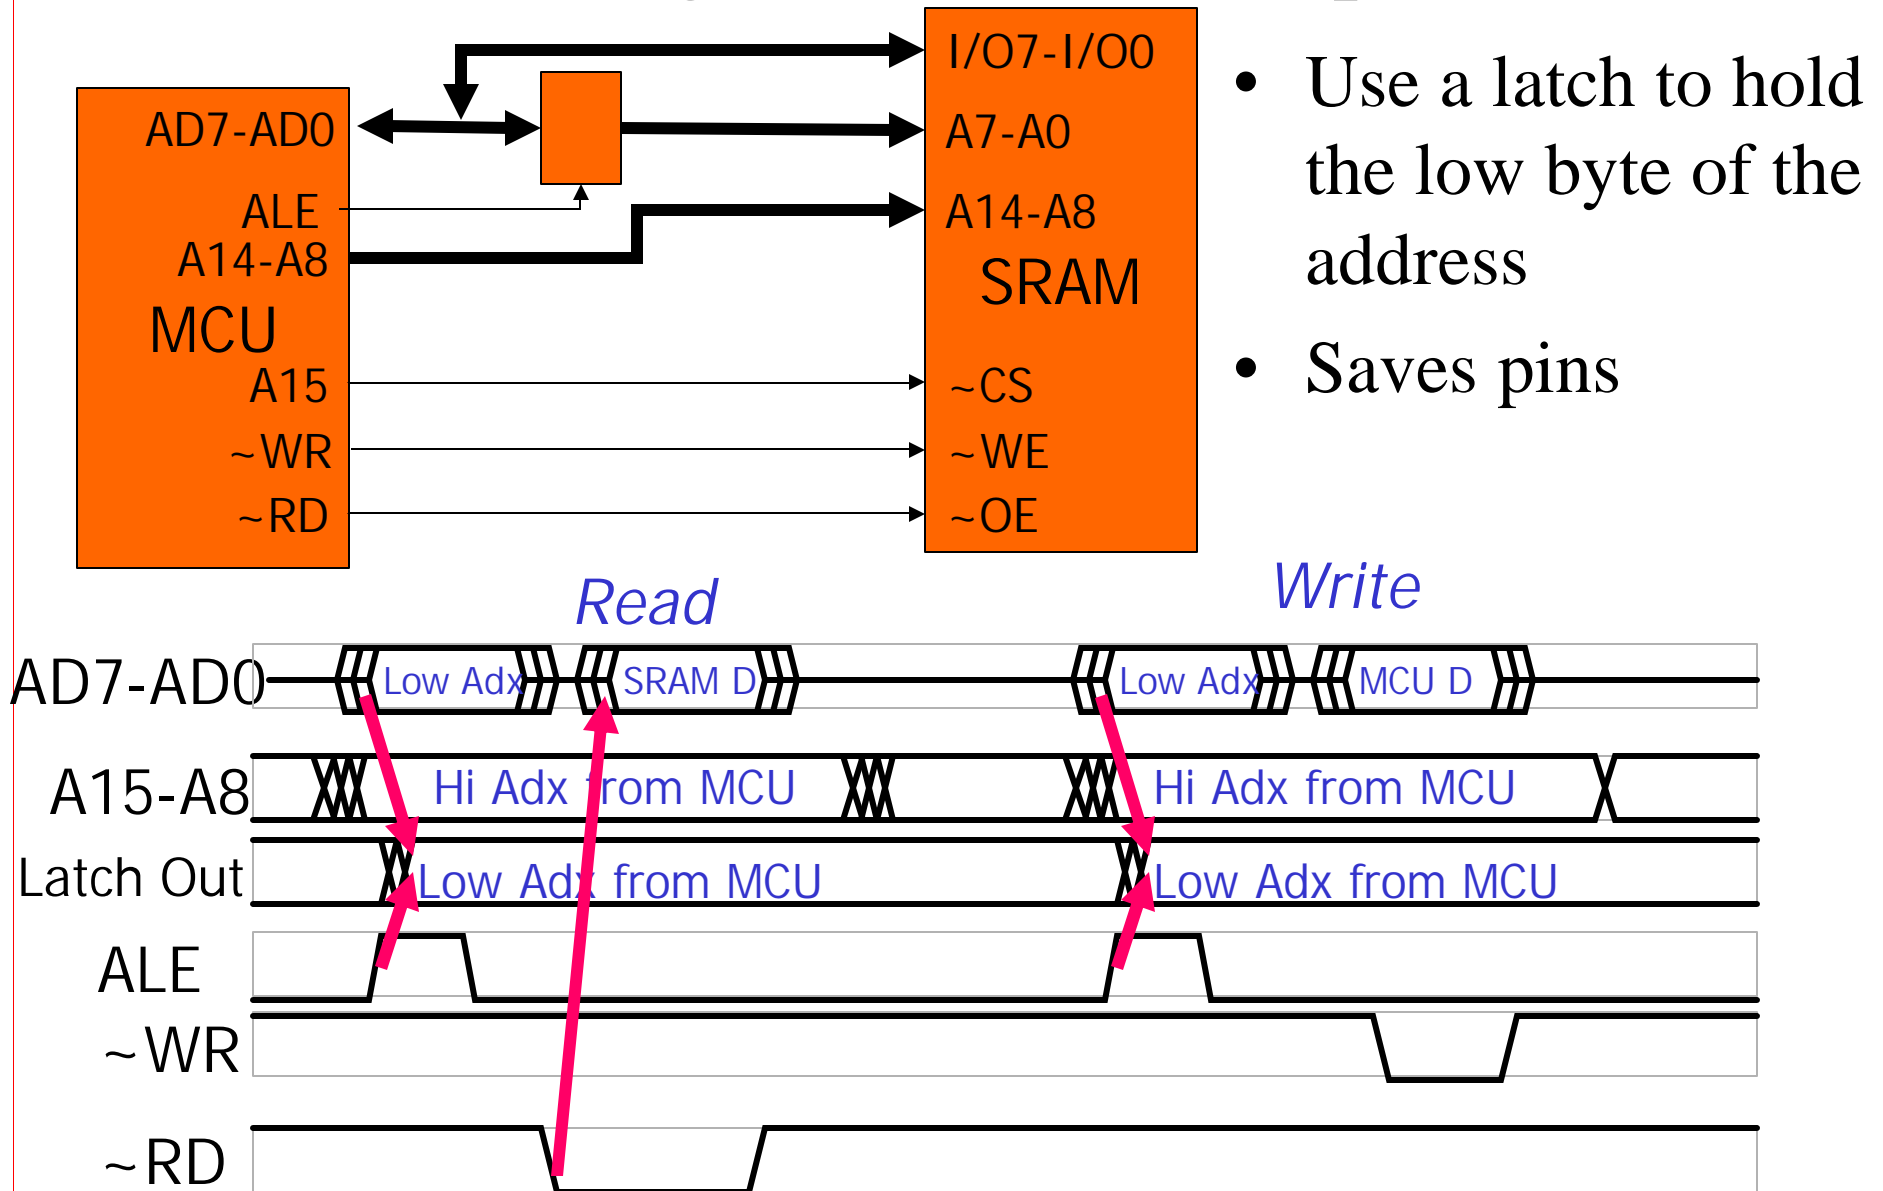
- Chip Select
 - Partial vs. Full Decoding
 - Power Consumption

Read

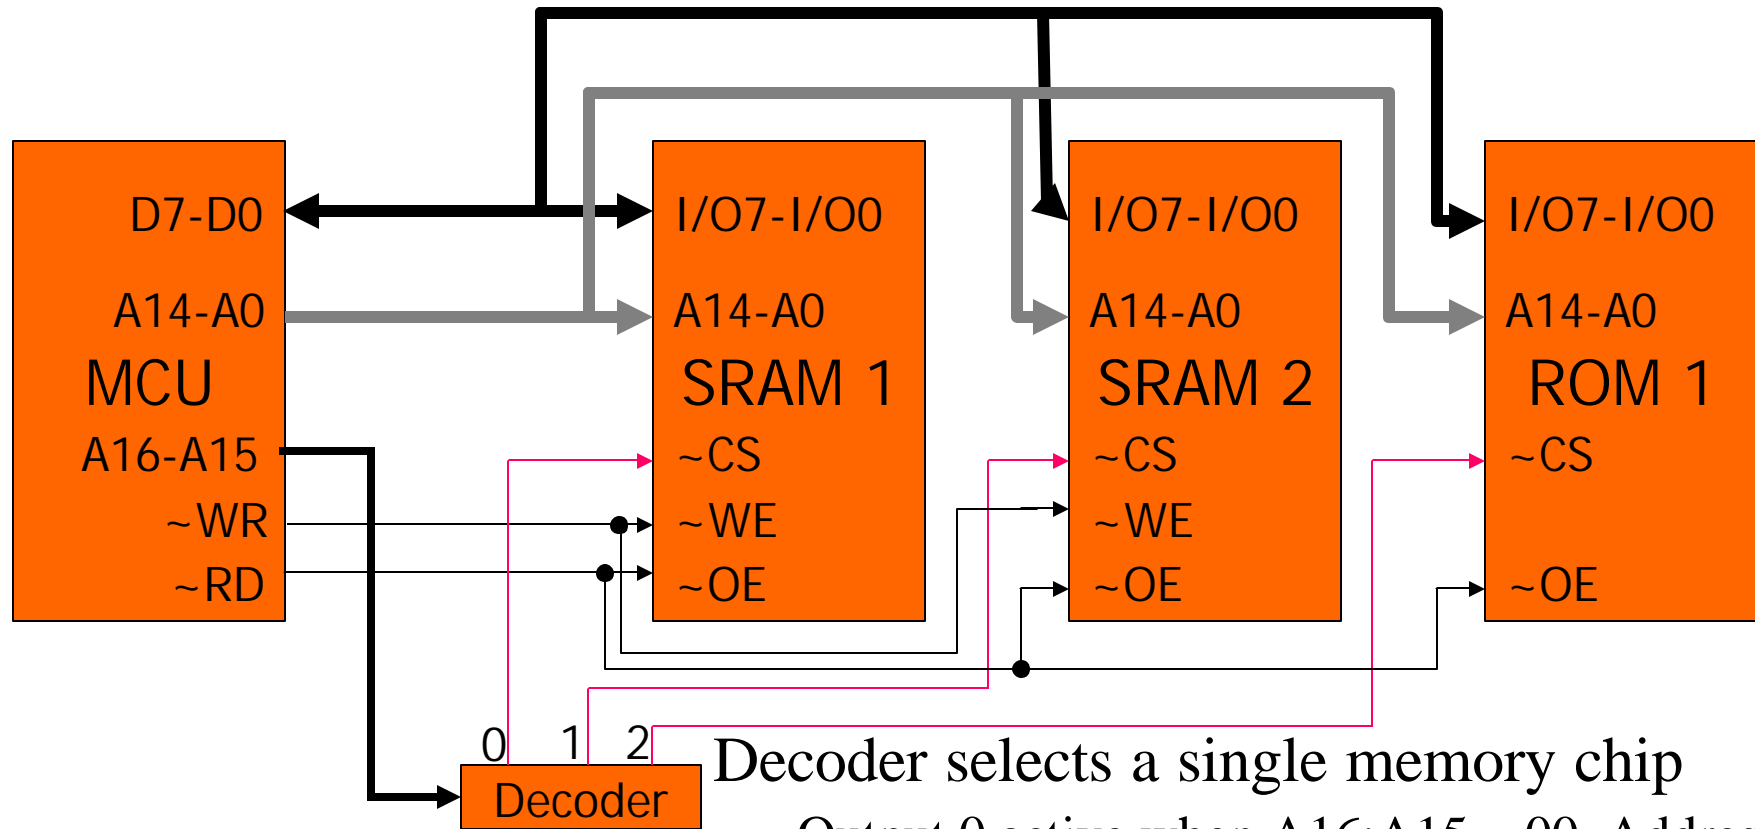
Write



External Memory Access – Multiplexed Buses



Multiple External Memories

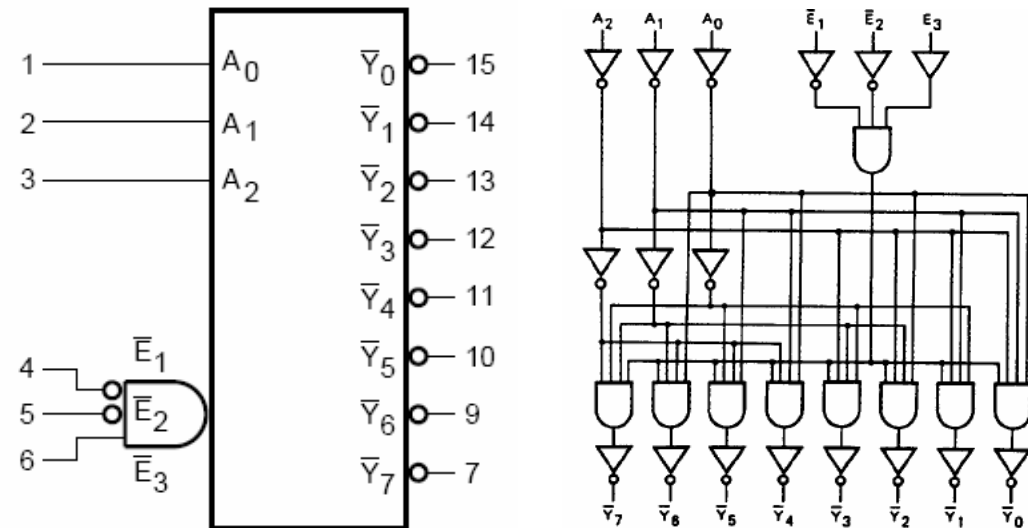


Decoder selects a single memory chip

- Output 0 active when A16:A15 = 00. Address = 0 0xxx xxxx xxxx xxxx = 00000h to 07FFFh
- Output 1 active when A16:A15 = 01. Address = 0 1xxx xxxx xxxx xxxx = 08000h to 0FFFFh
- Output 2 active when A16:A15 = 10. Address = 1 0xxx xxxx xxxx xxxx = 10000h to 17FFFh

Example: 74HC138 3-to-8 Line Decoder/Demultiplexer

- Common 74 series logic IC
- Diagrams from Philips Semiconductor
74HC/HCT138 Product Specification
- Typical propagation delay of up to 19 ns

[illegible]

Issue – Bus Loading

- Microcontroller has a limited output drive capacity on the data and address buses
 - e.g. 100 pF for an AMD MCU
- Each device on the bus adds capacitance
 - 11 pF for each input on IDT SRAM chip
- This leads to increased time delay until bus reaches valid voltage
 - M16C lowers threshold voltages when accessing external memory
- Solution: add buffers

Issue - The Memory Wall

- Difference in read cycle times for memories
 - Internal memories
 - SRAM: blazingly fast
 - Flash: down to 40 ns
 - External parallel-interface memories
 - SRAM: down to 10 ns
 - Flash: down to 55 ns
- Memory Wall: flash memory can't keep up with fast processors
 - External flash: $1/55 \text{ ns} = 18 \text{ MHz}$
 - Internal flash: $1/40 \text{ ns} = 25 \text{ MHz}$
- Need a mechanism to speed up access to data stored in flash memory



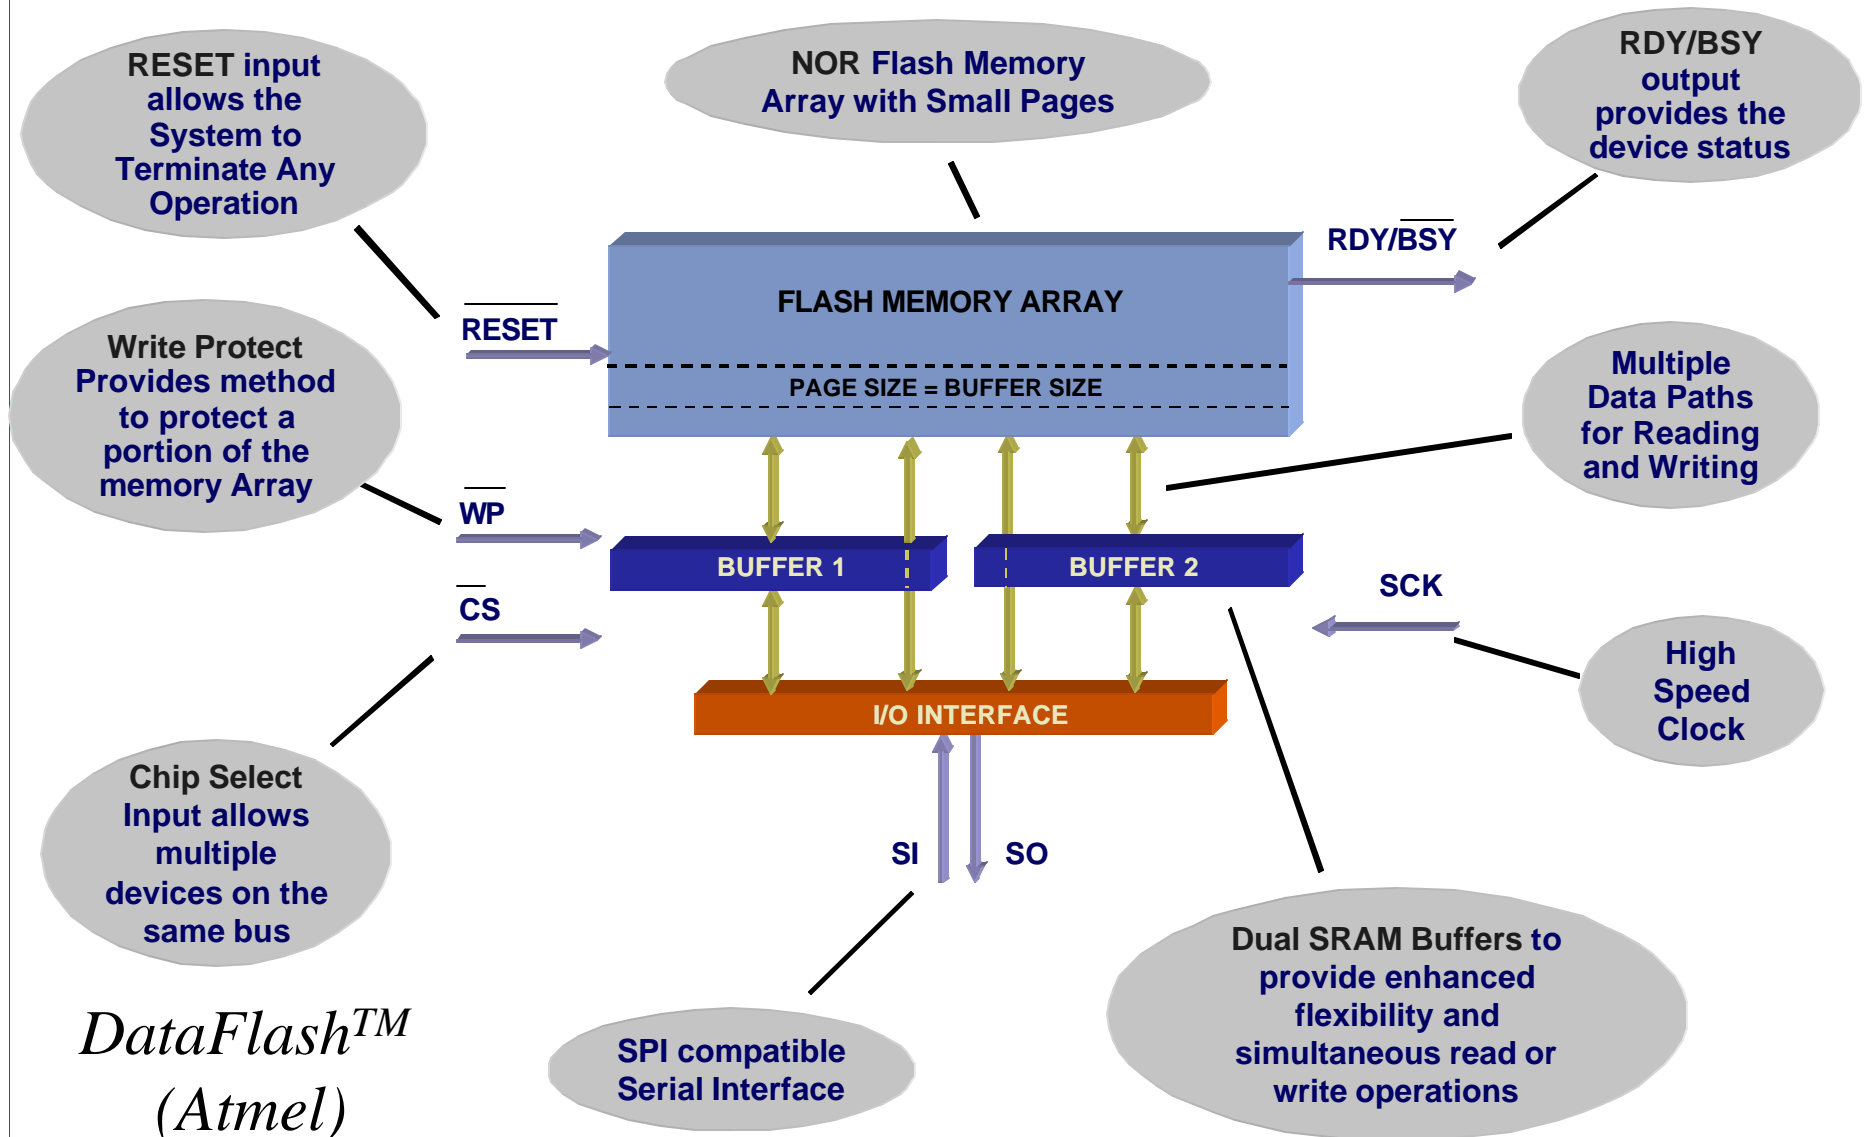
Solutions to the Embedded Memory Wall

- Code shadowing
 - Use flash memory to hold the program
 - At boot-time load the program into faster SRAM
 - Execute the program out of SRAM
 - Problem: extra memory costs money
- Cache
 - Use a small fast SRAM with a cache controller to hold commonly used data
 - Problem: unpredictable access times make it hard to guarantee predictable timing (e.g. real-time is difficult)
- Wider bus to memory
 - Access multiple (e.g. four bytes) at a time.
 - Fetching byte N also prefetches bytes N+1, N+2 and N+3
 - Works well for sequential accesses
 - Problem: Still incurs a delay for random accesses (e.g. branch)
- Branch target cache
 - Cache multiple bytes from targets of branch instructions
 - Problem: Unpredictable access times. But may be able to lock cache.
 - Could also cache the beginning of all ISRs

Storing Large Amounts of Data

- What if we want to store more data than fits into our processor's address space?
 - A. Use a memory paging scheme
 - Use a register (or output port) to hold the upper bits of the address
 - This register selects which *page* of memory we will access
 - Control the register with a page select function
 - This doesn't fit in well with C, as the compiler doesn't know about your custom paging scheme
 - Some MCUs support paging, and the compilers can compile for it
 - B. Use a serially-interfaced memory (inexpensive)
 - Communicate with the memory over a few data lines
 - C compiler doesn't know about this memory either
- May want to introduce some kind of *file system* to manage information in the expanded memory

Serial Interface Flash



Details for DataFlash

- Numbers
 - 1 Mbit – 128 Mbit capacities available
 - Page sizes from 264 to 1056 bytes
 - up to 20 MHz SPI interface speed
 - Page->Buffer transfer/compare: < 250 μ s
 - Page erase: < 8 ms
 - Page program: < 14 ms
- C source code for interface available
 - Flash file system (FAT12 and FAT16)
 - Compression/decompression
 - Error detection and correction
 - Wear leveling
 - Automatic page rewrite
- Commands (hardware)
 - Page read
 - Transfer page to buffer
 - Compare page with buffer
 - Buffer read, write
 - Program page from buffer (with or without erasing)
 - Program page through buffer
 - Page erase
 - Block erase (8 pages)
 - Automatic page rewrite

Moving Data Efficiently

- Sometimes we just need to move data...
 - Loading a packet from an Ethernet interface
 - Loading a video frame buffer
 - Initializing an array to zero
 - Loading an audio output buffer with audio samples

- Very slow when performed in software

```
Loop:    mov.w  _source[A0], _dest[A1]
         add.w  #2,  A0    ; increment src. ptr
         add.w  #2,  A1    ; increment dest. ptr
         cmp.w  A0,  R0    ; assume R0 is end ptr
         jle    Loop
```

- Consider a UART ISR which takes 50 cycles
 - Actually just need to move a byte from the UART to a buffer
 - $50 \text{ cycles} / 16 \text{ MHz} = 3.125 \text{ us}$
 - Limits us to maximum 320 kHz bit rate (100% CPU utilization)

Direct Memory Access

- Sequence of activities
 - Controller takes bus from CPU
 - Performs data transfer
 - Can interrupt CPU to signal completion
- Control Registers
 - Start and destination addresses
 - Transfer length
 - Status
- DMA Controller preempts CPU, may need to interleave accesses to ensure progress for CPU
- UART DMA: 2 cycles
 - At 320 kbps, CPU utilization = $320 \text{ kHz} * 2/16 \text{ MHz} * 100\% = 4\%$ utilization

M16C/62 DMA Controllers

- 2 DMACs, 0 and 1
- Control Registers
 - Source Address SAR0, SAR1
 - Destination Address DAR0, DAR1
 - Transfer Count TCR0, TCR1
 - DMA Request Cause
 - INT1, Timer, software trigger, UART Tx, UART Rx, A/D conversion complete
 - Transfer mode
 - Enable
 - Byte or word transfer
 - Single or repeat transfer
 - Source address increment or fixed
 - Destination address increment or fixed